

Po-han Li

pohanli@utexas.edu | [Personal Website](#) | [LinkedIn](#) | [Google Scholar](#)

RESEARCH INTERESTS

I build foundation models that turn messy, multimodal data into reliable predictions. My work spans tabular foundation models that predict instantly via in-context learning, cross-modal representation learning that translates video and sensor streams into language under tight data and compute budgets, and rigorous evaluation and human-AI systems that keep VLMs and LLMs trustworthy and aligned. I care about research that ships: state-of-the-art results that are efficient, open, and usable in production.

EDUCATION

University of Texas at Austin

Aug. 2023 – May 2026

Ph.D. in Electrical and Computer Engineering

Texas, U.S.A.

- Decision, Information, and Communications Engineering (DICE) track
- Co-advised by Prof. [Sandeep Chinchali](#) and Prof. [Ufuk Topcu](#)
- GPA: 3.93/4.00

University of Texas at Austin

Aug. 2021 – May 2023

M.S. in Electrical and Computer Engineering

Texas, U.S.A.

- Decision, Information, and Communications Engineering (DICE) track

National Taiwan University

Sep. 2016 – Jul. 2020

B.S. in Electrical Engineering

Taipei, Taiwan

- Research Advisor: Prof. Wanjiun Liao; GPA: 4.26/4.30, Ranking: 4/177
- Dean's List (thrice); College Student Research Scholarship, Ministry of Science and Technology (2017–2019)

WORK EXPERIENCE

SyntheFY

Jun. 2026 – Present

Founding Research Scientist

Austin, U.S.A (Remote)

- Develop [Nori](#), a tabular foundation model that predicts missing or future values in spreadsheets and SQL tables via in-context learning
- Build Nori for forecasting, pricing, risk scoring, and other structured-data prediction tasks with calibrated uncertainty and prediction intervals
- Led the design, pretraining, evaluation, and open-source release of Nori, which has passed 400K downloads across Hugging Face, PyPI, and GitHub
- Achieved **state-of-the-art regression accuracy** while keeping the model efficient and practical for real-world tabular prediction

Meta Platforms, Inc.

May 2025 – Aug. 2025

Software Engineer Intern @ Monetization GenAI

California, U.S.A

- Built automation tools to extract information from external websites and aligned source code with visual snapshots, supporting evaluation pipelines for ML-driven ad generation
- Engineered robust metrics to assess the effectiveness of unsupervised generated advertisements

Meta Platforms, Inc.

May 2024 – Aug. 2024

Software Engineer Intern @ Infra+Ranking & Foundational AI

California, U.S.A

- Calibrated Meta's multimodal foundation model for ads ranking and analyzed iterative-training performance to develop robust evaluation benchmarks
- Developed scalable classifiers for distributed clusters in Python, integrating heterogeneous data sources with preprocessing and exploratory analysis

China Network Systems Co., Ltd.

Oct. 2019 – Mar. 2021

Machine Learning and Data Scientist Intern

Taipei, Taiwan

- Built prediction models for **churn (unsubscribe) rate** and collected network-quality data via distributed Raspberry Pi
- Created databases and interactive reports to monitor **over 1M set-top boxes** in real-time

PUBLICATIONS

For a complete list of my publications, please check my [Google Scholar](#).

First-Author Publications:

1. **P. Li**, S. Chen, U. Topcu, and S. Chinchali. Visil: Unified evaluation of information loss in multimodal video captioning. *Under review*, 2026
2. S. Chen*, **P. Li***, X. Sun, S. Yang, E. Barsoum, Z. Liu, S. Chinchali, and U. Topcu. VEGAS: Human-aligned video caption evaluation via gaze. *Under review at Transactions on Machine Learning Research (TMLR)*, 2026
3. N. P. Bhatt*, **P. Li***, and K. Gupta* et. al. UNCAP: Uncertainty-guided planning using natural language communication for cooperative autonomous vehicles. In *International Conference on Autonomous Agents and Multi-Agent Systems (AAMAS) Best Paper Nomination*, 2026
4. **P. Li***, S. Chen*, S. Chinchali, and U. Topcu. VIBE: Video-to-text information bottleneck evaluation for TL;DR. *Advances in Neural Information Processing Systems (NeurIPS)*, 2025
5. **P. Li**, S. Chinchali, and U. Topcu. CSA: Data-efficient mapping of unimodal features to multimodal features. *International Conference on Learning Representations (ICLR)*, 2025
6. **P. Li**, S. K. Ankireddy, R. Zhao, H. N. Mahjoub, E. Moradi-Pari, U. Topcu, S. Chinchali, and H. Kim. Task-aware distributed source coding under dynamic bandwidth. *Advances in Neural Information Processing Systems (NeurIPS)*, 2023

Other Publications:

1. M. Omama, **P. Li**, and S. Chinchali. Exploiting distribution constraints for scalable and efficient image retrieval. *International Conference on Learning Representations (ICLR)*, 2025
2. O. Akcin, **P. Li**, S. Agarwal, and S. Chinchali. Decentralized data collection for robotic fleet learning: A game-theoretic approach. In *Conference on Robot Learning (CoRL)*, 2022
3. Y. Geng, D. Zhang, **P. Li**, O. Akcin, A. Tang, and S. Chinchali. Decentralized sharing and valuation of fleet robotic data. In *Conference on Robot Learning (CoRL)*, 2021

ACADEMIC SERVICE

Reviewer

Aug. 2021 - Present

- Reviewer of ICML, NeurIPS, ICLR, CVPR, MLSys, AAAI, AISTATS, IEEE Systems Journal, IROS, and ICRA.

EXTRACURRICULAR ACTIVITIES

UT Girl Day

Volunteer

Feb. 2023 and Feb. 2024

- Taught Python basics using AutoAuto cars; reached 200+ visitors of all ages. [link](#)

Code2College

Mentor

Jul. 2022 – Jul. 2024

- Mentored underrepresented high school graduates to prepare for software engineering jobs or college interviews.

REACT REU

Mentor

Aug. 2023

- Instructed undergraduate students to improve real-time computer vision models, excelling in image classification and object detection using Python.